

26-10-23

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```



```
In [5]: # Data mining
df = pd.read_csv('emp.csv')
df.head(1)
```

Out[5]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus	Senior Management	Team
0	Douglas	Male	08/06/1993	12:42 pm	97308	6.945	True	Marketing

```
In [6]: df.sample(5)
```

Out[6]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus	Senior Management	Team
247	Brenda	NaN	7/27/2010	11:07 pm	106115	3.742	True	Product
724	Andrea	Female	12/10/2001	6:40 am	37888	13.470	False	Engineering
361	Margaret	Female	05/05/2014	6:01 am	55044	4.078	False	Sales
290	Jeremy	Male	6/14/1988	6:20 pm	129460	13.657	True	NaN
776	Bobby	Male	9/30/2004	10:52 pm	79047	18.784	False	Human Resources

```
In [7]: df.isnull().sum()
```

```
Out[7]: First Name          67
Gender          145
Start Date       0
Last Login Time  0
Salary           0
Bonus           0
Senior Management 67
Team            43
dtype: int64
```

```
In [8]: df.isnull().all
```

```
Out[8]: <bound method NDFrame._add_numeric_operations.<locals>.all of      First Name
Gender  Start Date  Last Login Time  Salary  Bonus \
0      False  False      False      False  False  False  False
1      False  False      False      False  False  False  False
2      False  False      False      False  False  False  False
3      False  False      False      False  False  False  False
4      False  False      False      False  False  False  False
..      ...      ...      ...      ...      ...      ...
995     False  True      False      False  False  False  False
996     False  False     False      False  False  False  False
997     False  False     False      False  False  False  False
998     False  False     False      False  False  False  False
999     False  False     False      False  False  False  False

      Senior Management  Team
0              False  False
1              False   True
2              False  False
3              False  False
4              False  False
..              ...      ...
995             False  False
996             False  False
997             False  False
998             False  False
999             False  False

[1000 rows x 8 columns]>
```

```
In [9]: df.isnull().any()
```

```
Out[9]: First Name          True
Gender          True
Start Date      False
Last Login Time False
Salary          False
Bonus          False
Senior Management True
Team           True
dtype: bool
```

```
In [10]: df['Salary'].isnull()
```

```
Out[10]: 0      False
1      False
2      False
3      False
4      False
...
995    False
996    False
997    False
998    False
999    False
Name: Salary, Length: 1000, dtype: bool
```

```
In [11]: df['Salary'].isnull().all()
```

```
Out[11]: False
```

```
In [12]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   First Name            933 non-null   object
1   Gender                855 non-null   object
2   Start Date            1000 non-null  object
3   Last Login Time       1000 non-null  object
4   Salary                1000 non-null  int64
5   Bonus                 1000 non-null  float64
6   Senior Management     933 non-null   object
7   Team                  957 non-null   object
dtypes: float64(1), int64(1), object(6)
memory usage: 62.6+ KB
```

```
In [13]: df.describe()
```

```
Out[13]:
```

	Salary	Bonus
count	1000.000000	1000.000000
mean	90662.181000	10.207555
std	32923.693342	5.528481
min	35013.000000	1.015000
25%	62613.000000	5.401750
50%	90428.000000	9.838500
75%	118740.250000	14.838000
max	149908.000000	19.944000

In [14]: *# Data cleaning*

```
df['First Name']=df['First Name'].fillna('NA')  
df['First Name'].sample(25)
```

Out[14]:

13	Gary
741	Jane
347	Lori
152	Ruth
458	Albert
326	Jeffrey
456	Deborah
784	Stephanie
344	Scott
814	Rachel
928	Jeffrey
846	Stephen
376	Brandon
208	Jonathan
508	Scott
542	Amanda
278	Betty
320	NA
459	Charles
462	Craig
534	Gerald
989	Justin
623	Irene
686	Paul
762	Terry

Name: First Name, dtype: object

```
In [15]: df['First Name']=df['First Name'].replace(to_replace='NA', value='Seemi')
df['First Name'].sample(25)
```

```
Out[15]: 54      Sara
673      Ralph
555      Anne
319      Jacqueline
190      Carol
654      Carl
226      Kathy
462      Craig
894      Betty
491      Nicholas
576      Michael
676      Annie
930      Nancy
918      Ryan
832      Keith
970      Alice
460      Tina
675      Diane
218      Gregory
971      Patrick
90      Janice
823      Seemi
336      Mark
640      Kathleen
306      Mark
Name: First Name, dtype: object
```

```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
#   Column                Non-Null Count  Dtype
---  -
0   First Name            1000 non-null   object
1   Gender                855 non-null    object
2   Start Date            1000 non-null   object
3   Last Login Time       1000 non-null   object
4   Salary                1000 non-null   int64
5   Bonus                1000 non-null   float64
6   Senior Management     933 non-null    object
7   Team                  957 non-null    object
dtypes: float64(1), int64(1), object(6)
memory usage: 62.6+ KB
```

```
In [17]: df.dropna(inplace=True)
df
```

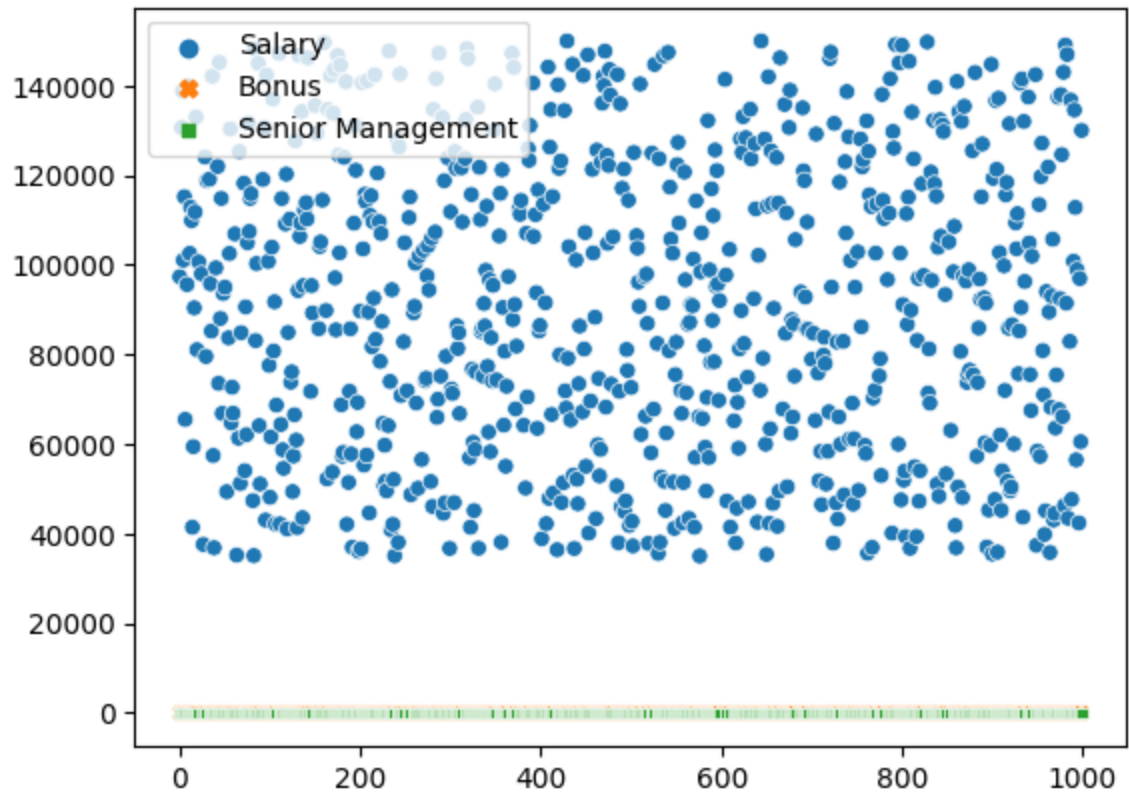
Out[17]:

	First Name	Gender	Start Date	Last Login Time	Salary	Bonus	Senior Management	Team
0	Douglas	Male	08/06/1993	12:42 pm	97308	6.945	True	Marketing
2	Maria	Female	4/23/1993	11:17 am	130590	11.858	False	Finance
3	Jerry	Male	03/04/2005	1:00 pm	138705	9.340	True	Finance
4	Larry	Male	1/24/1998	4:47 pm	101004	1.389	True	Client Services
5	Dennis	Male	4/18/1987	1:35 am	115163	10.125	False	Legal
...
994	George	Male	6/21/2013	5:47 pm	98874	4.479	True	Marketing
996	Phillip	Male	1/31/1984	6:30 am	42392	19.675	False	Finance
997	Russell	Male	5/20/2013	12:39 pm	96914	1.421	False	Product
998	Larry	Male	4/20/2013	4:45 pm	60500	11.985	False	Business Development
999	Albert	Male	5/15/2012	6:24 pm	129949	10.169	True	Sales

764 rows × 8 columns


```
In [18]: # sns.scatterplot(df['Salary'], df['Bonus%'])  
sns.scatterplot(df)
```

Out[18]: <Axes: >



```
In [19]: sns.scatterplot(x = 'Salary', y = 'Bonus', data = df, hue='Senior Management')
```

```
Out[19]: <Axes: xlabel='Salary', ylabel='Bonus'>
```



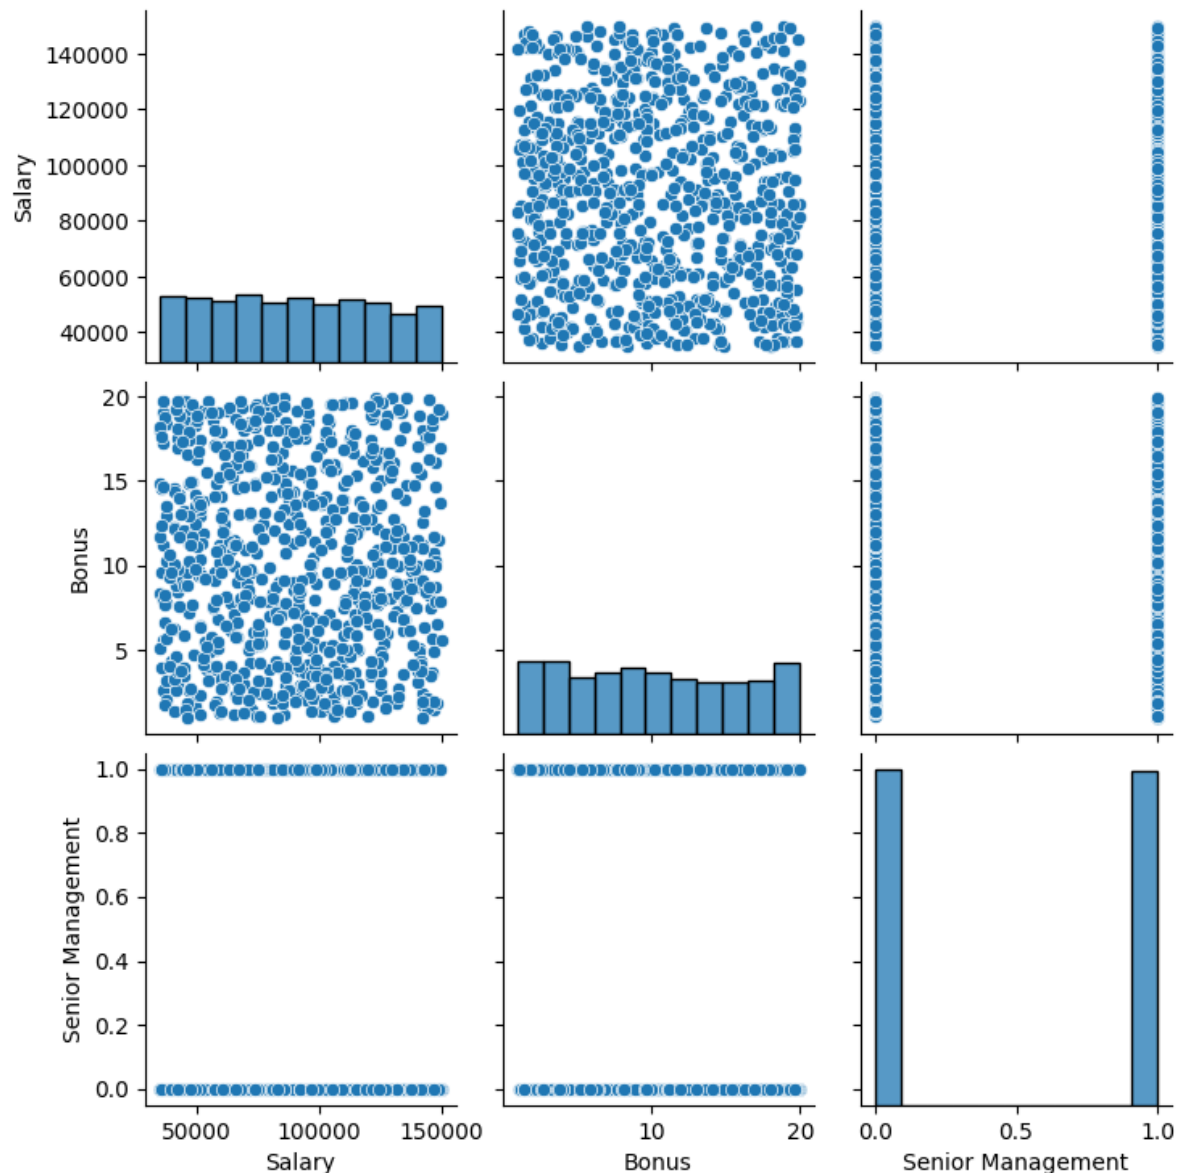

```
In [20]: sns.pairplot(df)
```

C:\Users\Administrator\AppData\Local\Programs\Python\Python311\Lib\site-packages\seaborn_stats\counting.py:137: RuntimeWarning: Converting input from bool to <class 'numpy.uint8'> for compatibility.

bin_edges = np.histogram_bin_edges(vals, bins, binrange, weight)
C:\Users\Administrator\AppData\Local\Programs\Python\Python311\Lib\site-packages\seaborn_stats\counting.py:176: RuntimeWarning: Converting input from bool to <class 'numpy.uint8'> for compatibility.

hist, edges = np.histogram(vals, **bin_kws, weights=weights, density=density)

```
Out[20]: <seaborn.axisgrid.PairGrid at 0x17fdbb03650>
```



```
In [21]: import ydata_profiling
```

```
In [22]: from ydata_profiling import ProfileReport
```

```
In [23]: profile = ProfileReport(df, title='Profiling Report')
```

```
In [24]: profile.to_file('report.html')
```

Summarize dataset:

100%

21/21 [00:02<00:00, 6.40it/s,

Completed]

Generate report structure:

100%

1/1 [00:03<00:00,

3.54s/it]

Render HTML: 100%

1/1 [00:01<00:00, 1.71s/it]

Export report to file:

100%

1/1 [00:00<00:00,

46.92it/s]

```
In [25]: profile.to_notebook_iframe()
```

Overview

Dataset statistics

Number of variables	8
Number of observations	764
Missing cells	0
Missing cells (%)	0.0%
Duplicate rows	0
Duplicate rows (%)	0.0%
Total size in memory	53.7 KiB
Average record size in memory	72.0 B

Variable types

Text	1
Categorical	2
DateTime	2
Numeric	2
Boolean	1

Reproduction

```
In [26]: # x =df.drop(['First Name', 'Gender', 'Start Date', 'Last Login', 'Bonus', 'Senior M  
x =df[['Salary']]
```

```
In [27]: y=df['Bonus']
```

```
In [28]: from sklearn.model_selection import train_test_split
```

```
In [29]: x_train, x_test, y_train, y_test = train_test_split(x,y)
```

```
In [30]: x_train.shape
```

```
Out[30]: (573, 1)
```

```
In [31]: y_test.shape
```

```
Out[31]: (191,)
```

```
In [32]: # Model Building
```

```
from sklearn.linear_model import LinearRegression  
model = LinearRegression()  
model
```

```
Out[32]: 

▼ LinearRegression



LinearRegression()


```

```
In [33]: # Training Model
```

```
model.fit(x_train, y_train)
```

```
Out[33]: 

▼ LinearRegression



LinearRegression()


```

```
In [34]: model.predict(x_test)
```

```
Out[34]: array([ 9.76089795, 10.2514      , 10.19062668,  9.88817301,  9.92210149,
 10.39675423, 10.36344088, 10.03094175, 10.51299532, 10.616655   ,
 10.1790737  ,  9.99366848, 10.66814364,  9.90909717, 10.69533448,
 10.48547769, 10.16151355, 10.29817901,  9.76754908, 10.03956324,
 10.01084417,  9.73642715, 10.53545732, 10.50695932, 10.18338925,
 10.16904894, 10.20646638, 10.74821678, 10.44212997, 10.34645741,
 10.48797667,  9.78468633, 10.57778623, 10.70949217, 10.62628569,
 10.49660777, 10.82737683, 10.34923513, 10.55656411,  9.98536417,
 10.33820116,  9.79242356, 10.1461352  , 10.15896651,  9.90901067,
 10.77217817, 10.77274524, 10.53472685, 10.61221451, 10.5501052  ,
 10.54903833, 10.38520124, 10.56047597, 10.81383427, 10.62530532,
 10.0738954  , 10.70175494, 10.06822464, 10.27985955, 10.29304649,
 10.30426307, 10.71486498, 10.46477459,  9.81703853, 10.77542684,
 10.48764988, 10.58136169, 10.60951368, 10.00159794, 10.07007965,
 10.8221578  , 10.00732638,  9.90488735, 10.52341415, 10.82234042,
 10.19123221, 10.33419318, 10.68645348, 10.54262748, 10.14726935,
  9.8283224  , 10.3066371  , 10.81283468, 10.44604183, 10.61855807,
 10.00847014, 10.12419222, 10.28736611, 10.0955116  ,  9.77263355,
 10.00350101,  9.79320209, 10.0914075  ,  9.83579051, 10.75999083,
 10.45510000, 10.40001000, 10.50001000,  9.00000000, 10.45510000,
```

```
In [35]: model.predict([[97308]])
```

```
C:\Users\Administrator\AppData\Local\Programs\Python\Python311\Lib\site-packa
ges\sklearn\base.py:465: UserWarning: X does not have valid feature names, bu
t LinearRegression was fitted with feature names
  warnings.warn(
```

```
Out[35]: array([10.24199037])
```

```
In [36]: y_pred= model.predict(x_test)
y_pred
```

```
Out[36]: array([ 9.76089795, 10.2514      , 10.19062668,  9.88817301,  9.92210149,
 10.39675423, 10.36344088, 10.03094175, 10.51299532, 10.616655   ,
 10.1790737  ,  9.99366848, 10.66814364,  9.90909717, 10.69533448,
 10.48547769, 10.16151355, 10.29817901,  9.76754908, 10.03956324,
 10.01084417,  9.73642715, 10.53545732, 10.50695932, 10.18338925,
 10.16904894, 10.20646638, 10.74821678, 10.44212997, 10.34645741,
 10.48797667,  9.78468633, 10.57778623, 10.70949217, 10.62628569,
 10.49660777, 10.82737683, 10.34923513, 10.55656411,  9.98536417,
 10.33820116,  9.79242356, 10.1461352  , 10.15896651,  9.90901067,
 10.77217817, 10.77274524, 10.53472685, 10.61221451, 10.5501052  ,
 10.54903833, 10.38520124, 10.56047597, 10.81383427, 10.62530532,
 10.0738954  , 10.70175494, 10.06822464, 10.27985955, 10.29304649,
 10.30426307, 10.71486498, 10.46477459,  9.81703853, 10.77542684,
 10.48764988, 10.58136169, 10.60951368, 10.00159794, 10.07007965,
 10.8221578  , 10.00732638,  9.90488735, 10.52341415, 10.82234042,
 10.19123221, 10.33419318, 10.68645348, 10.54262748, 10.14726935,
  9.8283224  , 10.3066371  , 10.81283468, 10.44604183, 10.61855807,
 10.00847014, 10.12419222, 10.28736611, 10.0955116  ,  9.77263355,
 10.00350101,  9.79320209, 10.0914075  ,  9.83579051, 10.75999083,
 10.45161649, 10.18881012, 10.52534606,  9.82387229, 10.15540065,
 10.10283553, 10.16746305, 10.3621145  , 10.15454523,  9.8287453  ,
 10.10677624, 10.55864018, 10.15391088,  9.7677317  , 10.04275425,
 10.05567206,  9.88790389, 10.00643251, 10.64820945,  9.85108235,
 10.40924913,  9.85678195,  9.95375205, 10.62103783, 10.46325597,
 10.2641448  ,  9.93101132, 10.69508458,  9.96323857, 10.83994863,
 10.60620734, 10.3052146  , 10.19771995,  9.81882627,  9.78016894,
 10.46884024, 10.59576928, 10.08079644,  9.83245533,  9.75366051,
 10.76639206, 10.39067017,  9.80844588, 10.2281018  , 10.6294767  ,
 10.29760232, 10.42419497, 10.39226567, 10.39859963, 10.35853903,
 10.53170885, 10.60318934,  9.98908381,  9.92511949, 10.31767107,
 10.14934543, 10.6019014  , 10.47665436, 10.45499011, 10.6832817  ,
 10.35204168, 10.21023408, 10.55090295,  9.97215802, 10.5980472  ,
 10.58042938,  9.98937215, 10.5435694  , 10.31758456, 10.23156193,
  9.89951454, 10.07809562, 10.70304288,  9.88285787, 10.03896733,
  9.85704146, 10.52210699,  9.78865587, 10.68536739, 10.68640543,
 10.66905673,  9.94659151, 10.14888408, 10.23143698, 10.82313817,
 10.78136673,  9.99222676, 10.07390502, 10.24847811, 10.16255159,
 10.48564108, 10.3143455  , 10.70148582, 10.68049437, 10.12471124,
  9.76527117])
```

```
In [37]: # Slop(m)
model.coef_
```

```
Out[37]: array([-9.61146802e-06])
```

```
In [38]: # intercept
model.intercept_
```

```
Out[38]: 11.17726309851513
```

```
In [39]: # y =mx+c
# x = 97308
# y =
# -8.56629539e-06*97308+10.926666573743816 =
```

```
In [40]: # verify model
from sklearn.metrics import r2_score, mean_squared_error
r2_score(y_test, y_pred)
```

Out[40]: -0.015701038209921858

```
In [41]: mean_squared_error(y_test, y_pred)
```

Out[41]: 33.78312178428922

```
In [42]: model.predict(df[['Salary']])
```

```
10.18227432, 10.64593153, 10.06511052, 10.08079644, 10.62430573,
 9.82138292, 10.31767107, 10.74821678, 10.11070733, 10.06822464,
10.39226567, 10.12396154, 10.28777794 , 9.80884956, 10.37623374,
10.01961944, 10.12419222, 10.42356061, 10.14934543, 10.33820116,
10.55656411, 10.68799132, 10.60318934, 10.68049437, 10.70109175,
10.56132178, 9.7580049 , 10.46663921, 10.78475958, 10.2700078 ,
10.77274524, 10.67669784, 10.84027542, 9.94879254, 10.81163324,
 9.96323857, 9.80704261, 10.49660777, 10.38160655, 10.16905855,
10.48623699, 10.11534005, 10.07081012, 10.70922305, 10.31936269,
10.3052146 , 10.21315596, 10.51278387, 10.69533448, 10.19746044,
10.63396526, 10.1850232 , 10.46477459, 10.45958439, 10.24024108,
10.17526756, 10.27046915, 10.6811864 , 10.15896651, 10.73436665,
 9.88346339, 10.1461352 , 9.81703853, 10.5435694 , 10.50456606,
 9.7626184 , 10.45500933, 9.9004853 , 10.74919715, 10.03606467,
10.72643719, 10.41236325, 9.98725763, 10.8244165 , 10.10283553,
10.48276725, 10.4921961 , 10.7261969 , 9.97371507, 10.0096908 ,
10.34696682, 10.39577386, 10.3621145 , 10.53545732, 10.00847014,
10.12471124, 9.98821878, 9.90488735, 9.75260325, 9.77410411,
10.63012067, 10.7790792 , 10.06561993, 10.43999622, 10.59684577,
10.74330532, 10.61254129, 10.44606106, 9.87857115, 10.82395515
```

```
In [43]: # Data Mining
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```



```
In [44]: df = pd.read_csv('car.csv')
df
```

Out[44]:

	User ID	Gender	Age	AnnualSalary	Purchased
0	385	Male	35	20000	0
1	681	Male	40	43500	0
2	353	Male	49	74000	0
3	895	Male	40	107500	1
4	661	Male	25	79000	0
...
995	863	Male	38	59000	0
996	800	Female	47	23500	0
997	407	Female	28	138500	1
998	299	Female	48	134000	1
999	687	Female	44	73500	0

1000 rows × 5 columns

```
In [45]: df.sample(5)
```

Out[45]:

	User ID	Gender	Age	AnnualSalary	Purchased
333	293	Male	28	89000	0
627	632	Female	18	86000	0
923	707	Female	59	102500	1
949	731	Male	46	80500	0
176	378	Female	53	82000	1

```
In [46]: # Data Cleaning
df.isnull().any()
```

```
Out[46]: User ID      False
Gender      False
Age         False
AnnualSalary False
Purchased   False
dtype: bool
```

```
In [47]: df['Gender'].unique()
```

```
Out[47]: array(['Male', 'Female'], dtype=object)
```


In [48]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   User ID         1000 non-null   int64
1   Gender          1000 non-null   object
2   Age             1000 non-null   int64
3   AnnualSalary    1000 non-null   int64
4   Purchased       1000 non-null   int64
dtypes: int64(4), object(1)
memory usage: 39.2+ KB
```

In [49]: `df.describe()`

Out[49]:

	User ID	Age	AnnualSalary	Purchased
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	500.500000	40.106000	72689.000000	0.402000
std	288.819436	10.707073	34488.341867	0.490547
min	1.000000	18.000000	15000.000000	0.000000
25%	250.750000	32.000000	46375.000000	0.000000
50%	500.500000	40.000000	72000.000000	0.000000
75%	750.250000	48.000000	90000.000000	1.000000
max	1000.000000	63.000000	152500.000000	1.000000

```
In [50]: df.drop('User ID', axis='columns', inplace=True)
df
```

Out[50]:

	Gender	Age	AnnualSalary	Purchased
0	Male	35	20000	0
1	Male	40	43500	0
2	Male	49	74000	0
3	Male	40	107500	1
4	Male	25	79000	0
...
995	Male	38	59000	0
996	Female	47	23500	0
997	Female	28	138500	1
998	Female	48	134000	1
999	Female	44	73500	0

1000 rows × 4 columns

```
In [51]: df['Gender'].replace({'Female':0, 'Male':1}, inplace=True)
df
```

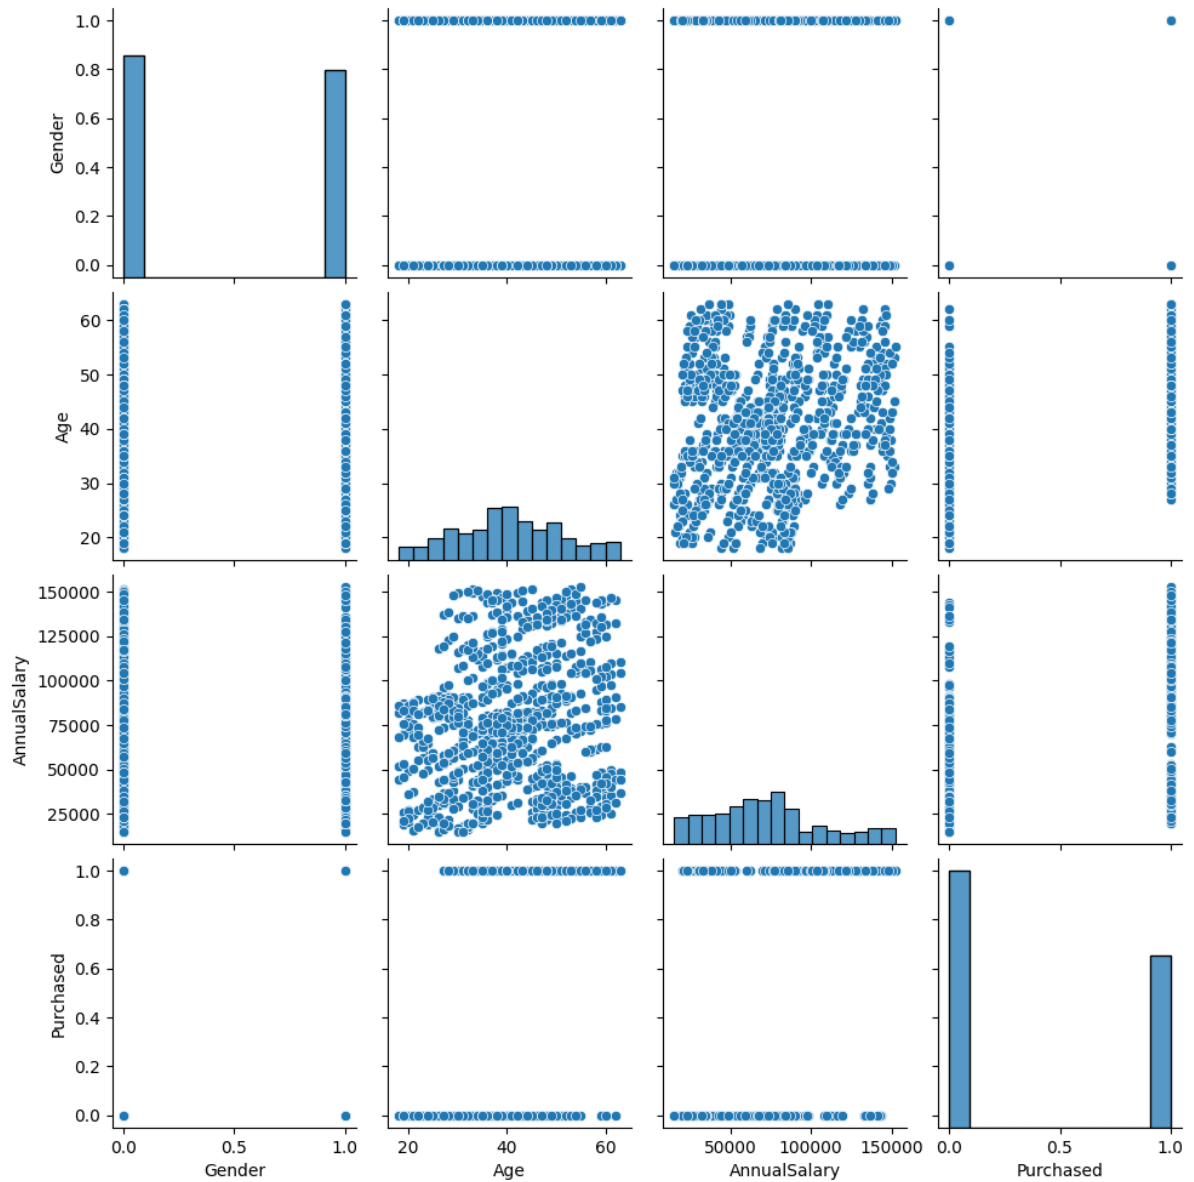
Out[51]:

	Gender	Age	AnnualSalary	Purchased
0	1	35	20000	0
1	1	40	43500	0
2	1	49	74000	0
3	1	40	107500	1
4	1	25	79000	0
...
995	1	38	59000	0
996	0	47	23500	0
997	0	28	138500	1
998	0	48	134000	1
999	0	44	73500	0

1000 rows × 4 columns

```
In [52]: sns.pairplot(df)
```

```
Out[52]: <seaborn.axisgrid.PairGrid at 0x17fe8ec3650>
```



```
In [53]: df.sample()
```

```
Out[53]:
```

	Gender	Age	AnnualSalary	Purchased
875	1	40	71000	1

```
In [54]: from sklearn.model_selection import train_test_split
```

```
In [55]: x_test, x_train, y_test, y_train = train_test_split (df[['Gender', 'Age', 'AnnualSalary', 'Purchased']], df['Purchased'], test_size=0.2, random_state=42)
```

```
In [56]: from sklearn.linear_model import LogisticRegression
```

```
In [57]: model = LogisticRegression()  
model
```

```
Out[57]: 

▼ LogisticRegression



LogisticRegression()


```

```
In [58]: model.fit(x_train, y_train)
```

```
Out[58]: 

▼ LogisticRegression



LogisticRegression()


```

```
In [59]: y_pred=model.predict(x_test)
y_pred
```

```
Out[59]: array([0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0, 1, 0,
 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1,
 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1,
 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0,
 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1,
 0, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 0,
 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 1, 0,
 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1,
 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 0,
 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1,
 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1,
 1, 0, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 1,
 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0,
 1, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1,
0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1,
1, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1,
0, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 1, 1, 1,
0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1,
1, 1, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 0,
1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0,
1, 0, 1, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0,
1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 0, 0,
1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1,
0, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1,
1, 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 1, 1,
1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0,
1, 0, 1, 0, 0, 1, 0, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1,
0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 1, 1, 1,
1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1,
0, 1, 0, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1,
1, 1, 0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1,
0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1,
1, 1], dtype=int64)
```

```
In [60]: model.score(x_test,y_test)
```

```
Out[60]: 0.372
```

```
In [61]: model.score(x_test, y_test)
```

```
Out[61]: 0.372
```

```
In [62]: model.coef_
```

```
Out[62]: array([[ -5.77313870e-04,  5.54036553e-03, -2.54997107e-06]])
```

In [63]: `model.intercept_`

Out[63]: `array([-0.00079754])`

In [64]: `x_train.sample()`

Out[64]:

	Gender	Age	AnnualSalary
818	0	38	79500

In [65]: `def sig(a):
import math
return 1/(1+math.e**(-a))`

```
b =0
c =45
d = 150000
round(pred_f(b,c,d))
```

In [66]: `def pred_f(b,c,d):
m = -7.20706471e-04*b+-9.59322438e-03*c+3.4273703e-06*d+(-0.00117166)
n =sig(m)
return n`

In [67]: `b =0
c =45
d = 150000
round(pred_f(b,c,d))`

Out[67]: 1

In [68]: `b = 0
c = 49
d =880000
round(pred_f(b,c,d))`

Out[68]: 1